

## Project Motivation

**Enzo’s struggles**

Enzo is a community-developed adaptive mesh refinement simulation code designed for rich multi-physics hydrodynamic astrophysical calculations. Unfortunately, Enzo’s scalability is limited by its design and implementation.

- Memory usage**
  - AMR structure is not scalable
  - permanent 3-layer ghost zones
  - memory fragmentation
- Data locality**
  - disrupted by load balancing
- Mesh quality**
  - 2-to-1 refinement violated
- Parallel task definition**
  - varying patch sizes
  - determined by AMR
- Parallel task scheduling**
  - parallel within a level
  - synchronization between levels

**Enzo-P/Cello’s solutions**

Enzo-P is being developed to bring Enzo’s powerful multi-physics capabilities to petascale and beyond. Our approach is to build on **Cello**, a scalable forest-of-octree AMR framework implemented using **Charm++**.

- Memory usage**
  - AMR structure fully distributed
  - ghosts allocated on demand
  - uniformity reduces fragmentation
- Data locality**
  - nearest-neighbor communication
- Mesh quality**
  - 2-to-1 refinement maintained
- Parallel task definition**
  - uniform sized blocks
  - user-specified size
- Parallel task scheduling**
  - asynchronous, data-driven
  - block-local time stepping

## Cello Components

Cello is organized into layered components for representing scientific problems, operating on fully-distributed data-structures, and interfacing with hardware, the user, and (later) other applications.

**High-level**

- Simulation:** represents a numerical simulation
- Problem:** defines a multiphysics problem
- Method:** implements a computational method

**Data structures**

- Mesh:** represents an AMR mesh
- Field:** represents Eulerian field data
- Particle:** represents Lagrangian particle data

**Middle-level**

- Control:** handles sequencing and synchronization
- Parameters:** reads and manages input file parameters
- Io:** coordinates disk data output

**Hardware interface**

- Disk:** interfaces to file format libraries
- Memory:** manages dynamic memory allocation
- Parallel:** interface to parallel library (deprecated)

**User interface**

- Portal:** interacts with other applications
- Monitor:** informs user of simulation progress

**Cross-cutting**

- Performance:** collects performance data
- Error:** detects and evaluates anomalies

## Parallelization

**Charm++**

Cello is implemented using **Charm++**, a C++-based OOP parallel programming system developed independently at the University of Illinois, Urbana-Champaign. Its design tenets include efficient portability, latency tolerance, dynamic load balancing, and code reuse.

- Charm++ programs
  - Charm++ objects are *chares*
  - remote callable *entry methods*
  - communicate via *messages*
- Charm++ runtime system
  - maps chares to processors
  - schedules entry methods
  - load balances chares
  - supports checkpoint / restart

**Cello AMR**

Cello implements a fully-distributed forest-of-octree mesh hierarchy using a Charm++ *chare array* named *Block*. Blocks are indexed using a bit coding of the array (forest) and octree indices. Charm++ caches array indices for high performance and scalability.

- Asynchronous execution
  - Field data sent when available
  - computes when all ghosts received
- Dynamic scheduling
  - multiple Blocks per process
  - automatic overlap of comp./comm.
- Charm++ provides
  - dynamic load balancing
  - checkpoint / restart

## Project Development

Summary	Roadmap	Status																					
<p>This project serves as a transition from single-developer to distributed community open development.</p> <p>Development is pipelined, with staged Cello framework capabilities enabling new Enzo-P physics.</p> <p><b>This material is based upon work supported by the National Science Foundation under Grant No. SI2-SSE-1440709.</b></p>	<table border="1"> <thead> <tr> <th>Code release</th> <th>Cello development</th> <th>Enzo-P development</th> </tr> </thead> <tbody> <tr> <td>V1.0 (<i>hydro</i>)</td> <td>Elliptic problems</td> <td>Chemistry</td> </tr> <tr> <td>V2.0 (<i>chemistry</i>)</td> <td>Particle support</td> <td>Self-gravity / RHD</td> </tr> <tr> <td>V3.0 (<i>gravity</i>)</td> <td>MHD support</td> <td>P<sup>3</sup>M / TreePM</td> </tr> <tr> <td>V4.0 (<i>particles</i>)</td> <td>Ray support</td> <td>MHD</td> </tr> <tr> <td>V5.0 (<i>magnetism</i>)</td> <td>optimization</td> <td>Adaptive ray tracing</td> </tr> <tr> <td>V6.0 (<i>radiation</i>)</td> <td></td> <td></td> </tr> </tbody> </table>	Code release	Cello development	Enzo-P development	V1.0 ( <i>hydro</i> )	Elliptic problems	Chemistry	V2.0 ( <i>chemistry</i> )	Particle support	Self-gravity / RHD	V3.0 ( <i>gravity</i> )	MHD support	P <sup>3</sup> M / TreePM	V4.0 ( <i>particles</i> )	Ray support	MHD	V5.0 ( <i>magnetism</i> )	optimization	Adaptive ray tracing	V6.0 ( <i>radiation</i> )			<p>Enzo-P / Cello is available under the BSD New Open Source license.</p> <ul style="list-style-type: none"> <li>Previous capabilities <ul style="list-style-type: none"> <li>PPM dual-energy hydrodynamics</li> <li>PPML compressible ideal MHD</li> </ul> </li> <li>Recent capabilities <ul style="list-style-type: none"> <li>self-gravity (Krylov)</li> <li>particles (static AMR)</li> </ul> </li> <li>Upcoming capabilities <ul style="list-style-type: none"> <li>self-gravity (MG, HG, P<sup>3</sup>M)</li> <li>particles (dynamic AMR)</li> <li>chemistry (Grackle library)</li> </ul> </li> </ul>
Code release	Cello development	Enzo-P development																					
V1.0 ( <i>hydro</i> )	Elliptic problems	Chemistry																					
V2.0 ( <i>chemistry</i> )	Particle support	Self-gravity / RHD																					
V3.0 ( <i>gravity</i> )	MHD support	P <sup>3</sup> M / TreePM																					
V4.0 ( <i>particles</i> )	Ray support	MHD																					
V5.0 ( <i>magnetism</i> )	optimization	Adaptive ray tracing																					
V6.0 ( <i>radiation</i> )																							

## Enzo-P Scaling

Preliminary scaling tests on Blue Waters indicate that Cello AMR can scale extremely well, especially in terms of memory.

**Test Problem**

- 3D “array” of Sedov blasts
- One blast per Blue Waters FP core
- 16<sup>3</sup> grid zones per block
- 10M blocks in 3-level AMR mesh
- Impossible with Enzo: **35GB** / process!

**Parallel Scaling**

- Efficiency at **P = 32K** FP cores:
  - parallel: **ε<sub>P</sub> = 0.868**
  - memory: **ε<sub>M</sub> = 0.993**

## Enzo-P/Cello Classes

Cello’s classes allow Enzo-P and other Cello application developers to easily extend their application’s functionality through inheritance.

**Cello classes**

- Simulation:** Simulation “chare group”
- Problem:** Problem definition
- Block:** Mesh block “chare array”

**Cello interfaces**

- Initial:** Initial conditions
- Boundary:** Boundary conditions
- Refine:** Refinement criteria
- Method:** Physics methods
- Output:** Disk output

**Enzo-P implementations**

- InitialFoo:** Enzo-P problem initializers
- BoundaryFoo:** Boundary condition types
- RefineFoo:** Specific refinement criteria
- MethodFoo:** Computational kernels
- OutputFoo:** Data or image output

## Cello Applications

**1. Parameter file: heat.in**

```
% list field variables
Field { list = [ "temperature" ]; }

% define initial conditions
Initial {
  list = [ "value" ];
  value {
    temperature = [ 10.0,
                    (0.3<x && x<0.7) &&
                    (0.3<y && y<0.7),
                    1.0 ];
  }
}

% define method parameters
Method {
  list = [ "heat" ];
  heat {
    alpha = 1.0;
  }
}
```

**Developing with Cello**

Writing scalable AMR applications with Cello is straightforward. For example, to use Cello to solve the heat equation with the Forward Euler method, two main steps are required:

- Create an input parameter file
- Add a new Method class

Other minor modifications are also needed for reading method parameters, calling the method’s constructor, and updating a Charm++ control file.

**2a. Include file: MethodHeat.hpp**

```
class MethodHeat : public Method {
  // Create a MethodHeat object
  MethodHeat ( double a ) :a.(a){};
  // Apply FE to the block
  virtual void compute(Block *);
  // Compute the CFL restriction
  virtual double timestep (Block *);
}
```

**2b. Source code: MethodHeat.cpp**

```
MethodHeat::compute(Block * block) {
  // Get field block attributes
  Field field = block()->data()->field();
  iU = field.field_id ("temperature");
  U = (double *) field.values (iU);
  field.dimensions (iU,&mx,&my);
  field.size (&nx,&ny);
  field.ghosts (iU,&gx,&gy);
  field.cell_width (&hx,&hy);
  rx = 1.0/(hx*hx); ry = 1.0/(hy*hy);
  // Apply forward Euler method
  for (int iy=gy; iy<ny+gy; iy++) {
    for (int ix=gx; ix<nx+gx; ix++) {
      i = ix + mx*iy;
      Uxx=(U[i-dx]-2*U[i]+U[i+dx])*rx;
      Uyy=(U[i-dy]-2*U[i]+U[i+dy])*ry;
      Unew[i]=U[i] + a_*dt*(Uxx + Uyy);
    } } }
```

